

FROM ZERO TO 80

YOUR FIRST NODE.JS WEB APPLICATION



Andrew Theken



Andrew Theken
@atheken
andrewtheke.com

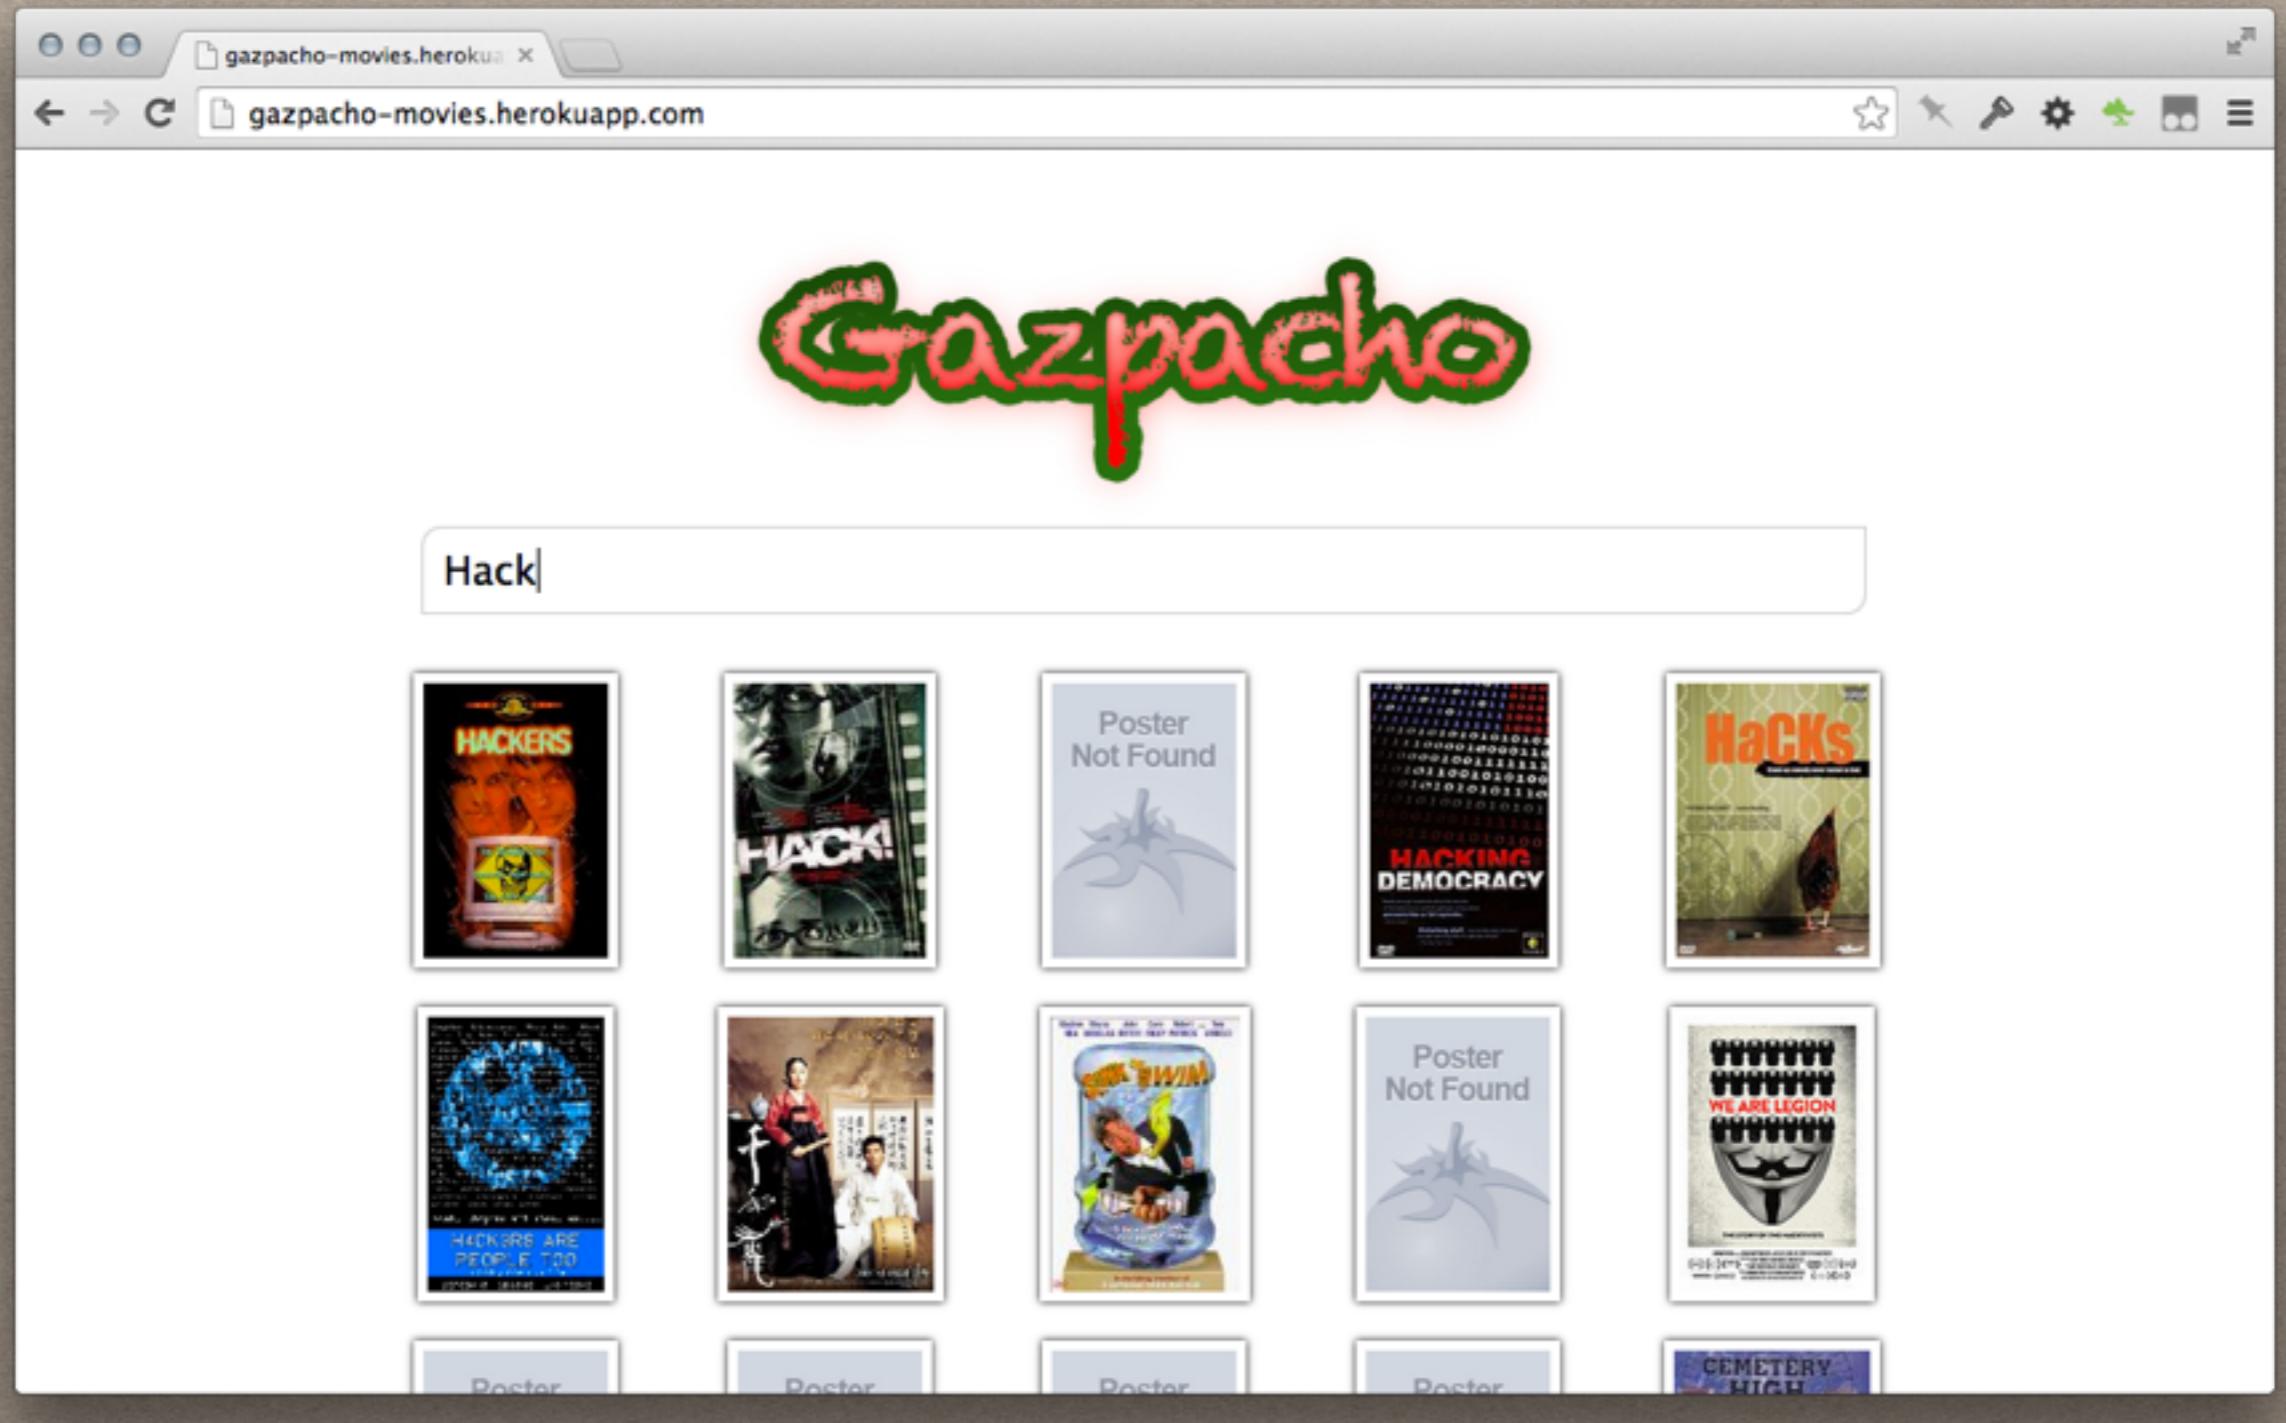


Curtis Herbert

WHAT WE'LL COVER

- What's Node?
- The reasons to use (or not use) Node.js
- Getting Node.js installed, and getting comfortable in the environment.
- Modules
- Events
- Streams
- Creating an HTTP server
- Middleware
 - Node.js “express”
 - HTML templating using Jade.

WHAT WE'LL BUILD



WHAT'S NODE.JS?

Node.js is a platform built on Chrome's JavaScript runtime for easily building fast, scalable network applications. Node.js uses an event-driven, non-blocking I/O model that makes it lightweight and efficient, perfect for data-intensive real-time applications that run across distributed devices.

- nodejs.org

WHY SHOULD YOU USE NODE.JS?

- One language everywhere: Javascript.
- Tons of modules available via NPM.
- Easy to deploy/run
- More than just HTTP Servers: CLI-apps, networked services, embedded systems, desktop GUI apps (really!)
- Active, fantastic community.
- It's Fun!

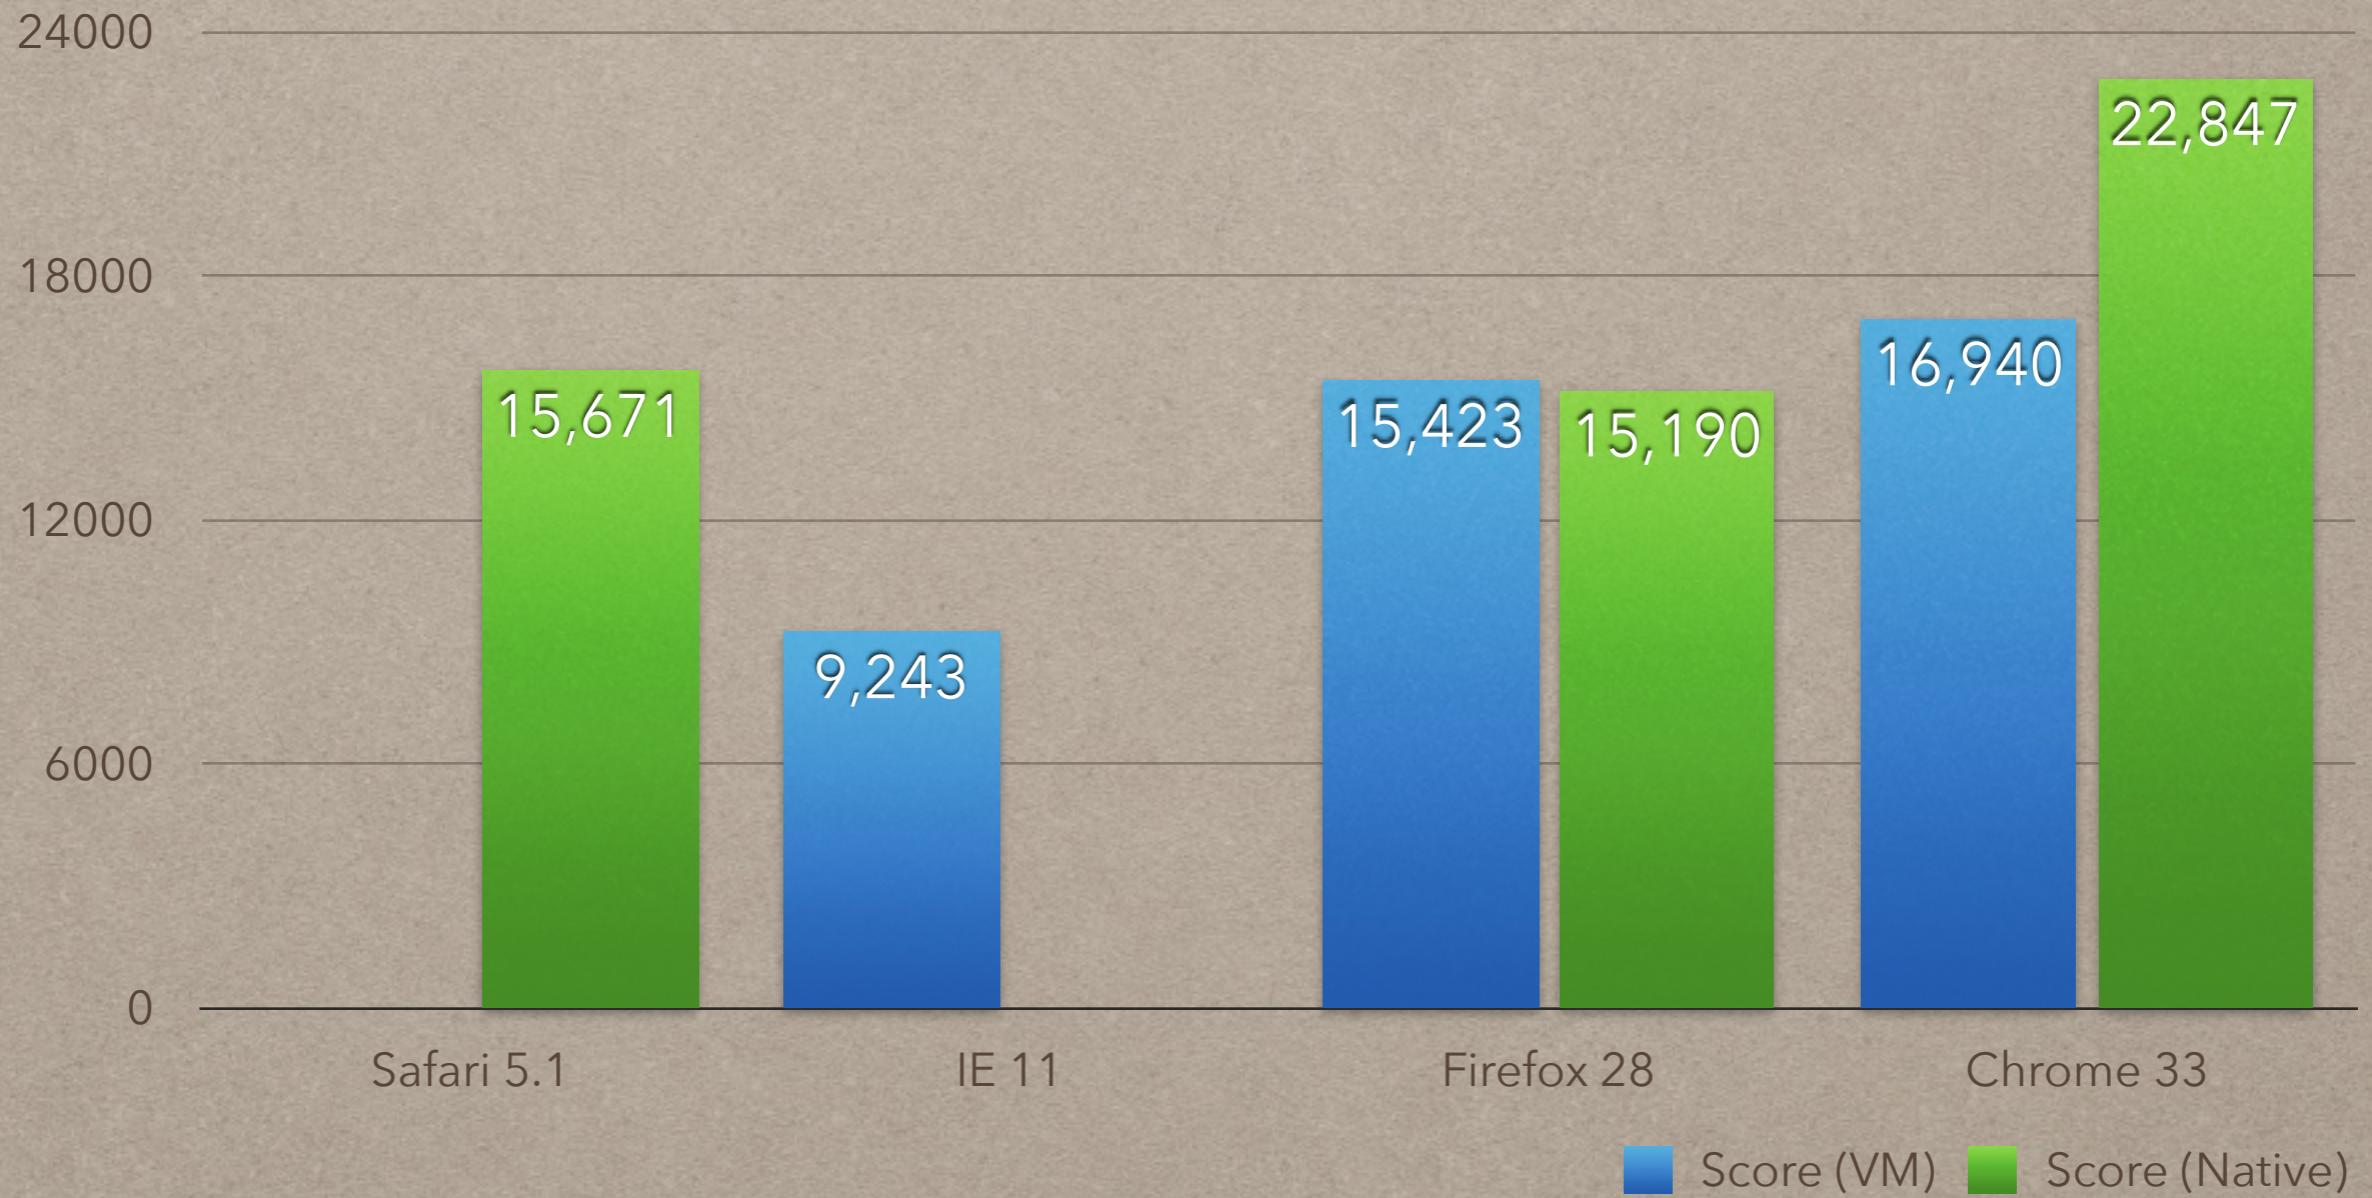
WHY SHOULDN'T YOU USE NODE.JS?

- Your app is CPU-intensive (graphics processing, for example)*
- Your app needs high-precision math capabilities.
- You expect to have a large-scale codebase (50kLoC+). This can be done, but requires practice and discipline.

* One could farm out CPU intensive work to a C++ process; "Think async."

WHY V8?

Chrome's V8 Javascript Engine is *FAST*.



octane-benchmark.googlecode.com/svn/latest/index.html

Exercise 1: Getting to know node

- Install Node.js
- Run the Node.js Prompt
- Print “hello world”
- Create your first node script.
- Explore the “globals” available in the node console.

<http://bit.ly/1mRtVHu>

NPM <http://npmjs.org>

The screenshot shows the homepage of npmjs.org. At the top left is the 'npm' logo. The top right features a search bar labeled 'Search Packages' and a user icon with 'Create Account | Login'. On the left, there's a sidebar with links to 'HOME', 'API', 'BLOG', 'NODE.JS', and 'JOBS'. The main content area has a large red 'npm' logo at the top. Below it is the title 'Node Packaged Modules'. It displays download statistics: 'Total Packages: 66 942', '8 340 364 downloads in the last day', '48 859 203 downloads in the last week', and '186 211 955 downloads in the last month'. There are also sections for 'Patches welcome!', 'Any package can be installed by using `npm install`', and 'Add your programs to this index by using `npm publish`'. At the bottom, there are two sections: 'Recently Updated' and 'Most Depended Upon', each listing several packages with their names and download counts.

Recently Updated	Most Depended Upon
• 3m stripify	• 5509 underscore
• 4m dockerode-process	• 4698 async
• 6m bugtrace	• 3898 request
• 6m dotject	• 2648 commander
• 9m docker-dns	• 2645 express
• 9m gulp-plates	• 2549 lodash
• 11m addressit	• 2530 optimist
• 11m vun-programmer	• 2111 coffee-script

Exercise 2: Doing something useful(?)

```
→ exercise2 git:(master) ✘ node colorize.js  
The quick brown fox jumps over the lazy dog.  
→ exercise2 git:(master) ✘ █
```

- Search npmjs.org for “ansi-colorizer”
- Review docs on “ansi-colorizer”
- install “ansi-colorizer” using NPM.
- check out the “node_modules” directory
- use “require” to load the ansi-colorizer
- Colorize vowels in a string.

EVENT-DRIVEN?

React to changes rather than “poll” for them.

*Imperative (**Active**):*

```
while(true){  
    if(isNewFileInDirectory("./uploads")){  
        processFile();  
    }  
}
```

*Declarative (**Passive**):*

```
var fileWatcher = new FileWatcher("./uploads");  
  
fileWatcher.on('newfile', processFile);
```

NON-BLOCKING I/O

I/O Is **slow**. Like, really, *really*, slow.

Reading **1Mb** from traditional hard drive takes 60,000,000 CPU Cycles

Event	Latency	Scaled
1 CPU cycle	0.3 ns	1 s
Level 1 cache access	0.9 ns	3 s
Level 2 cache access	2.8 ns	9 s
Level 3 cache access	12.9 ns	43 s
Main memory access (DRAM, from CPU)	120 ns	6 min
Solid-state disk I/O (flash memory)	50–150 µs	2–6 days
Rotational disk I/O	1–10 ms	1–12 months
Internet: San Francisco to New York	40 ms	4 years
Internet: San Francisco to United Kingdom	81 ms	8 years
Internet: San Francisco to Australia	183 ms	19 years
TCP packet retransmit	1–3 s	105–317 years
OS virtualization system reboot	4 s	423 years
SCSI command time-out	30 s	3 millennia
Hardware (HW) virtualization system reboot	40 s	4 millennia
Physical system reboot	5 m	32 millennia

Table 2.2 Example Time Scale of System Latencies

From: *Systems Performance: Enterprise and the Cloud*, p. 20

NON-BLOCKING I/O

Blocking I/O (**Synchronous, Active**):

```
console.log('Dictionary data loading. . .');
//Open and read a 1MB file from disk.
var content = fileSystem.readFileSync('./dictionary.txt');

//“1 year” later. . .
console.log('Dictionary loaded.');
```

Non-Blocking I/O (**Asynchronous, Passive**):

```
//Send request for file.
fileSystem.readFile('./dictionary.txt', function(err, data){
  if(!err){
    console.log('Dictionary loaded.');
  }
});

//CPU is ready for more work immediately.
console.log('Dictionary data loading...');
```

SYNC VS. ASYNC



Synchronous

"I'll wait here while you get this done"



Asynchronous

"Give me a job, get in, get out."

HOW DOES “ASYNC” WORK?

JAVASCRIPT THREAD

```
var fs = require('fs');
function cb(err, data){...};
fs.readFile('file.txt', cb);
```

BACKGROUND WORKER
THREADS (C++)

```
retrieveFile('file.txt');
```



HOW DOES “ASYNC” WORK?

JAVASCRIPT THREAD

```
var fs = require('fs');
function cb(err, data){...};
fs.readFile('file.txt', cb);
console.log('waiting...');
```

BACKGROUND WORKER
THREADS (C++)

```
retrieveFile('file.txt');
read bytes from FS
```



HOW DOES “ASYNC” WORK?

JAVASCRIPT THREAD

```
var fs = require('fs');
function cb(err, data){...};
fs.readFile('file.txt', cb);
console.log('waiting...');
cb(err,bytes);
```

BACKGROUND WORKER
THREADS (C++)

```
retrieveFile('file.txt');
read bytes from FS
return bytes
```



BIG DATA?

Congratulations! You work for wikipedia.

Your job is to find articles that could be enhanced by the inclusion of a cute cat pictures like the one on the right:



STREAMS

You decide that any article that mentions *cats*, *kittens*, or *felis* should include a kitten picture at the top.

There are **4,485,774** articles in the English version of Wikipedia. The full text of Wikipedia is **multiple terabytes**. Loading all of these entries into memory at once is not an option.

But, you're smart, and a node.js expert, so you recognize you can use *streams* to incrementally process the ginormous database of articles, one at a time.

Exercise 3: Data River

- `require('fs')`, check out the node docs for `FileSystem`
- Open utf-8 read file stream for “`dictionary.txt`”, and listen for “`data`” events.
- Colorize vowels in the file and write them to `process.stdout` as data becomes available.

FOLLOWUP: KITTEN OVERLOAD

There have 1,000,000,000+ edits to Wikipedia.

How can we efficiently track & update new or modified articles that might benefit from kittenification?



That's it for the "core" node concepts.

Let's do something fun!

HTTP WITH NODE

- Node's web server is built-in!
- An HTTP “transaction” is an input *stream* (request), and an output *stream* (response):
 - Requests have headers, parameters, urls, bodies, and verbs.
 - Responses have headers, bodies, and statuses.

THE (ALMOST) SIMPLEST WEB SERVER POSSIBLE

```
var http = require('http');

// create a server respond to all requests
// with the current date.
var server = http.createServer(
  function(request, response){
    response.write(new Date().toString());
    response.end();
});

//start the server on port 4514
server.listen(4514);
```

Exercise 4: Time and Temperature (ok, just time)

- Create a basic web server.
- Respond to any request with the current date and time on port 4514
- Open <http://localhost:4514/>
- Enhance the server to respond with message of ‘Not Found. Sorry.’ for any url that isn’t to ‘/datetime’

COMMON HTTP SERVER TASKS

- Logging requests
- Serve “static” files
- Compress responses
- Store posted files
- Parse request parameters and cookies
- Authenticate/
Authorize Requests
- Send non-2xx responses (404, 302, 500, etc.)
- Set cache headers
- Serve dynamic requests

Break for Lunch!

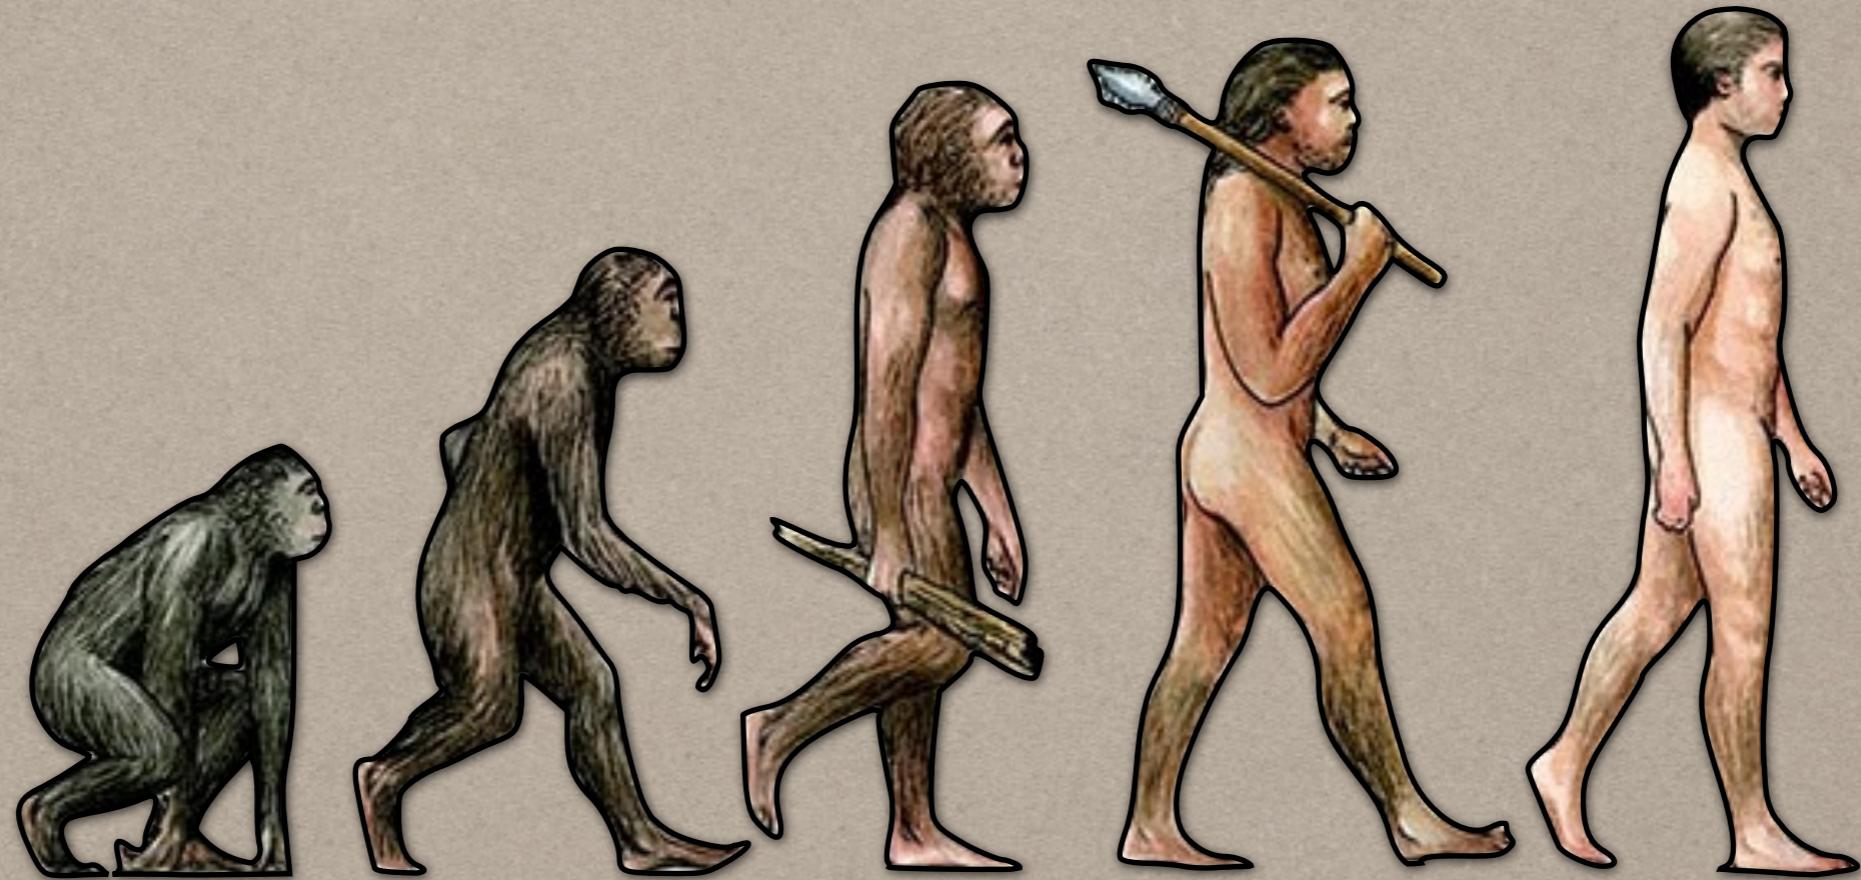
REVIEWING A SIMPLE HTTP SERVER

```
var http = require('http');

// create a server respond to all requests
// with the current date.
var server = http.createServer(
  function(request, response){
    if(request.url == '/datetime'){
      response.write(new Date().toString());
    }else{
      response.write('Not Found. Sorry.');
    }
    response.end();
});

//start the server on port 4514
server.listen(4514);
```

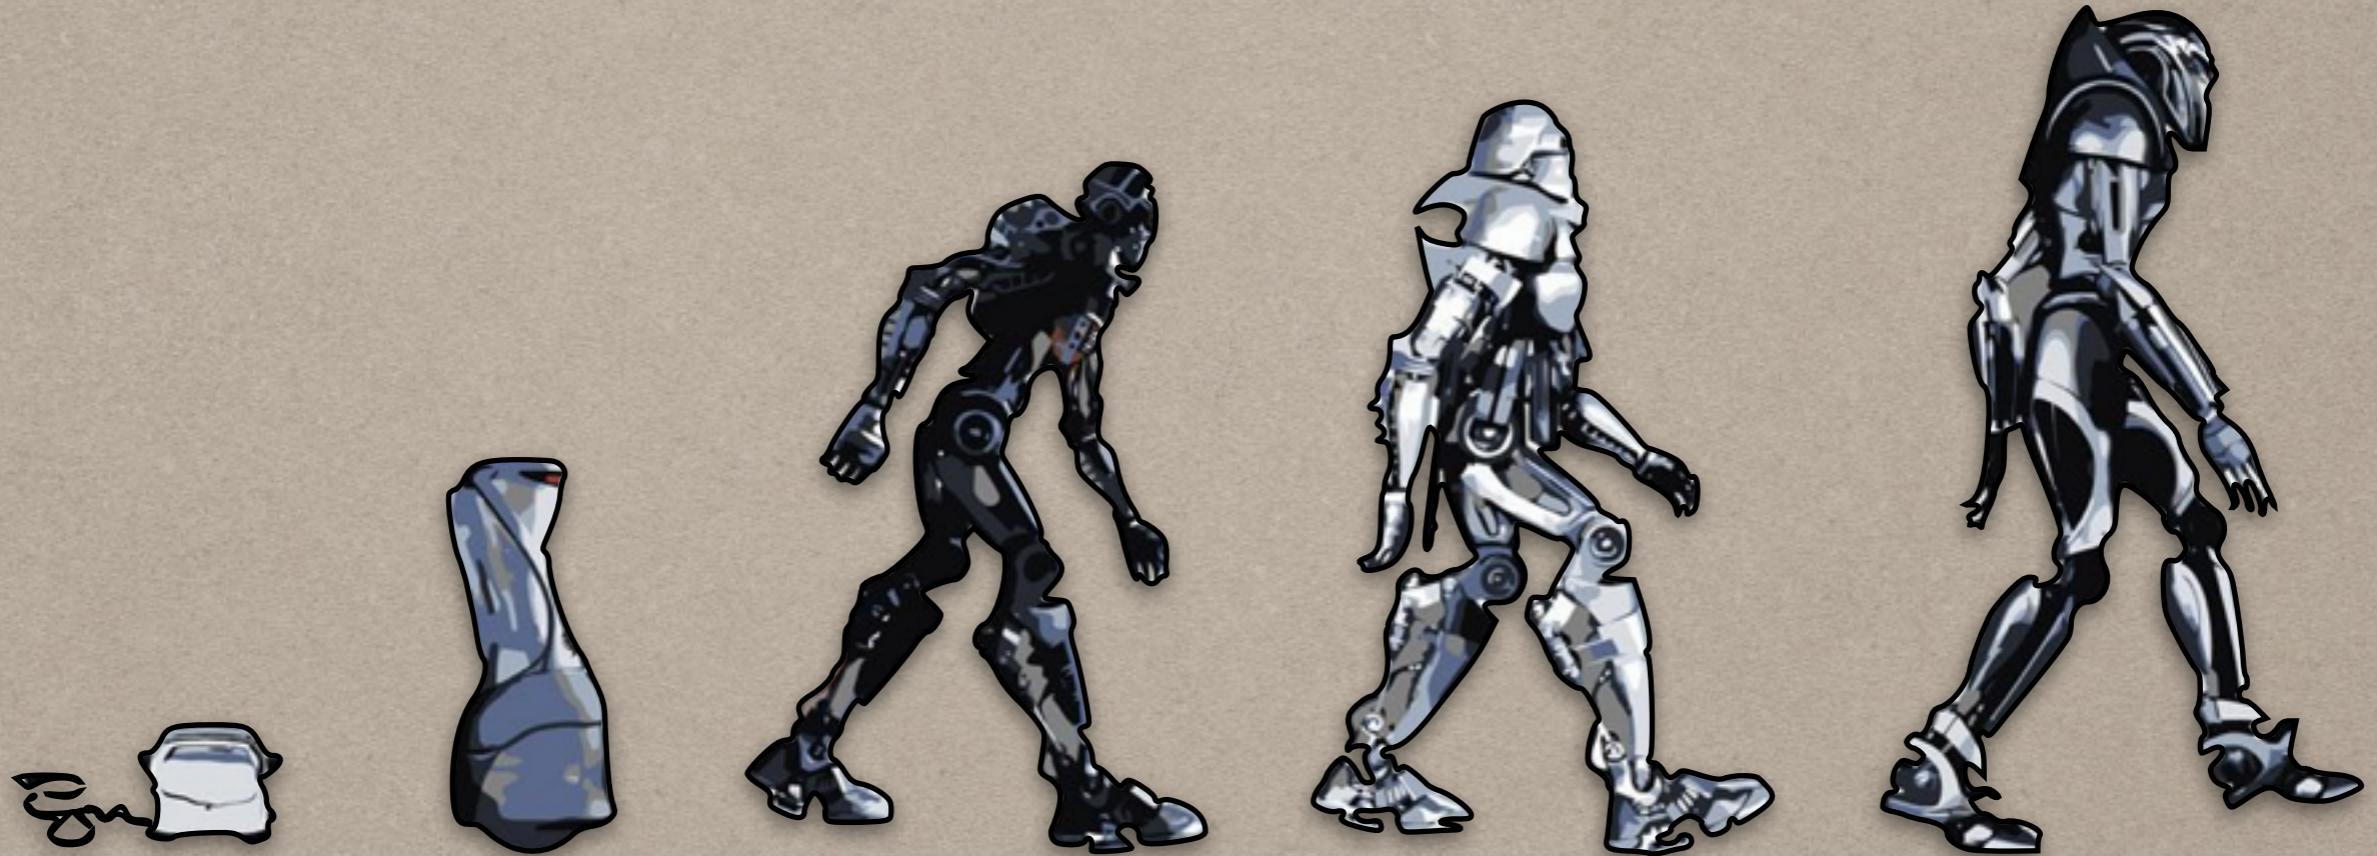
MIDDLEWARE



“Middleware is so simple,
you'd almost overlook explaining it.”

- me, just now

MIDDLEWARE



"Middleware is so simple,
you'd almost overlook explaining it."

- me, just now

THE MIDDLEWARE FUNCTION

Middleware Prototype:

```
function(request, response, done){  
    //inspect request  
    //manipulate response  
    //or, do nothing  
}
```

THE MIDDLEWARE FUNCTION

Middleware Prototype:

```
function(request, response, done){  
    //inspect request  
    //manipulate response  
    //or, do nothing  
}
```

Log Request:

```
//A simple request logger.  
function(request, response, done){  
    console.log(request.path);  
    done();  
}
```

THE MIDDLEWARE FUNCTION

Log Request:

```
//A simple request logger.  
function(request, response, done){  
  console.log(request.path);  
  done();  
}
```

Send content:

```
//A simple date-time response  
function(req, res, done){  
  if(req.url === '/datetime'){  
    res.write(new Date());  
    res.end();  
  }  
  else{  
    done();  
  }  
}
```

COMMON HTTP SERVER TASKS

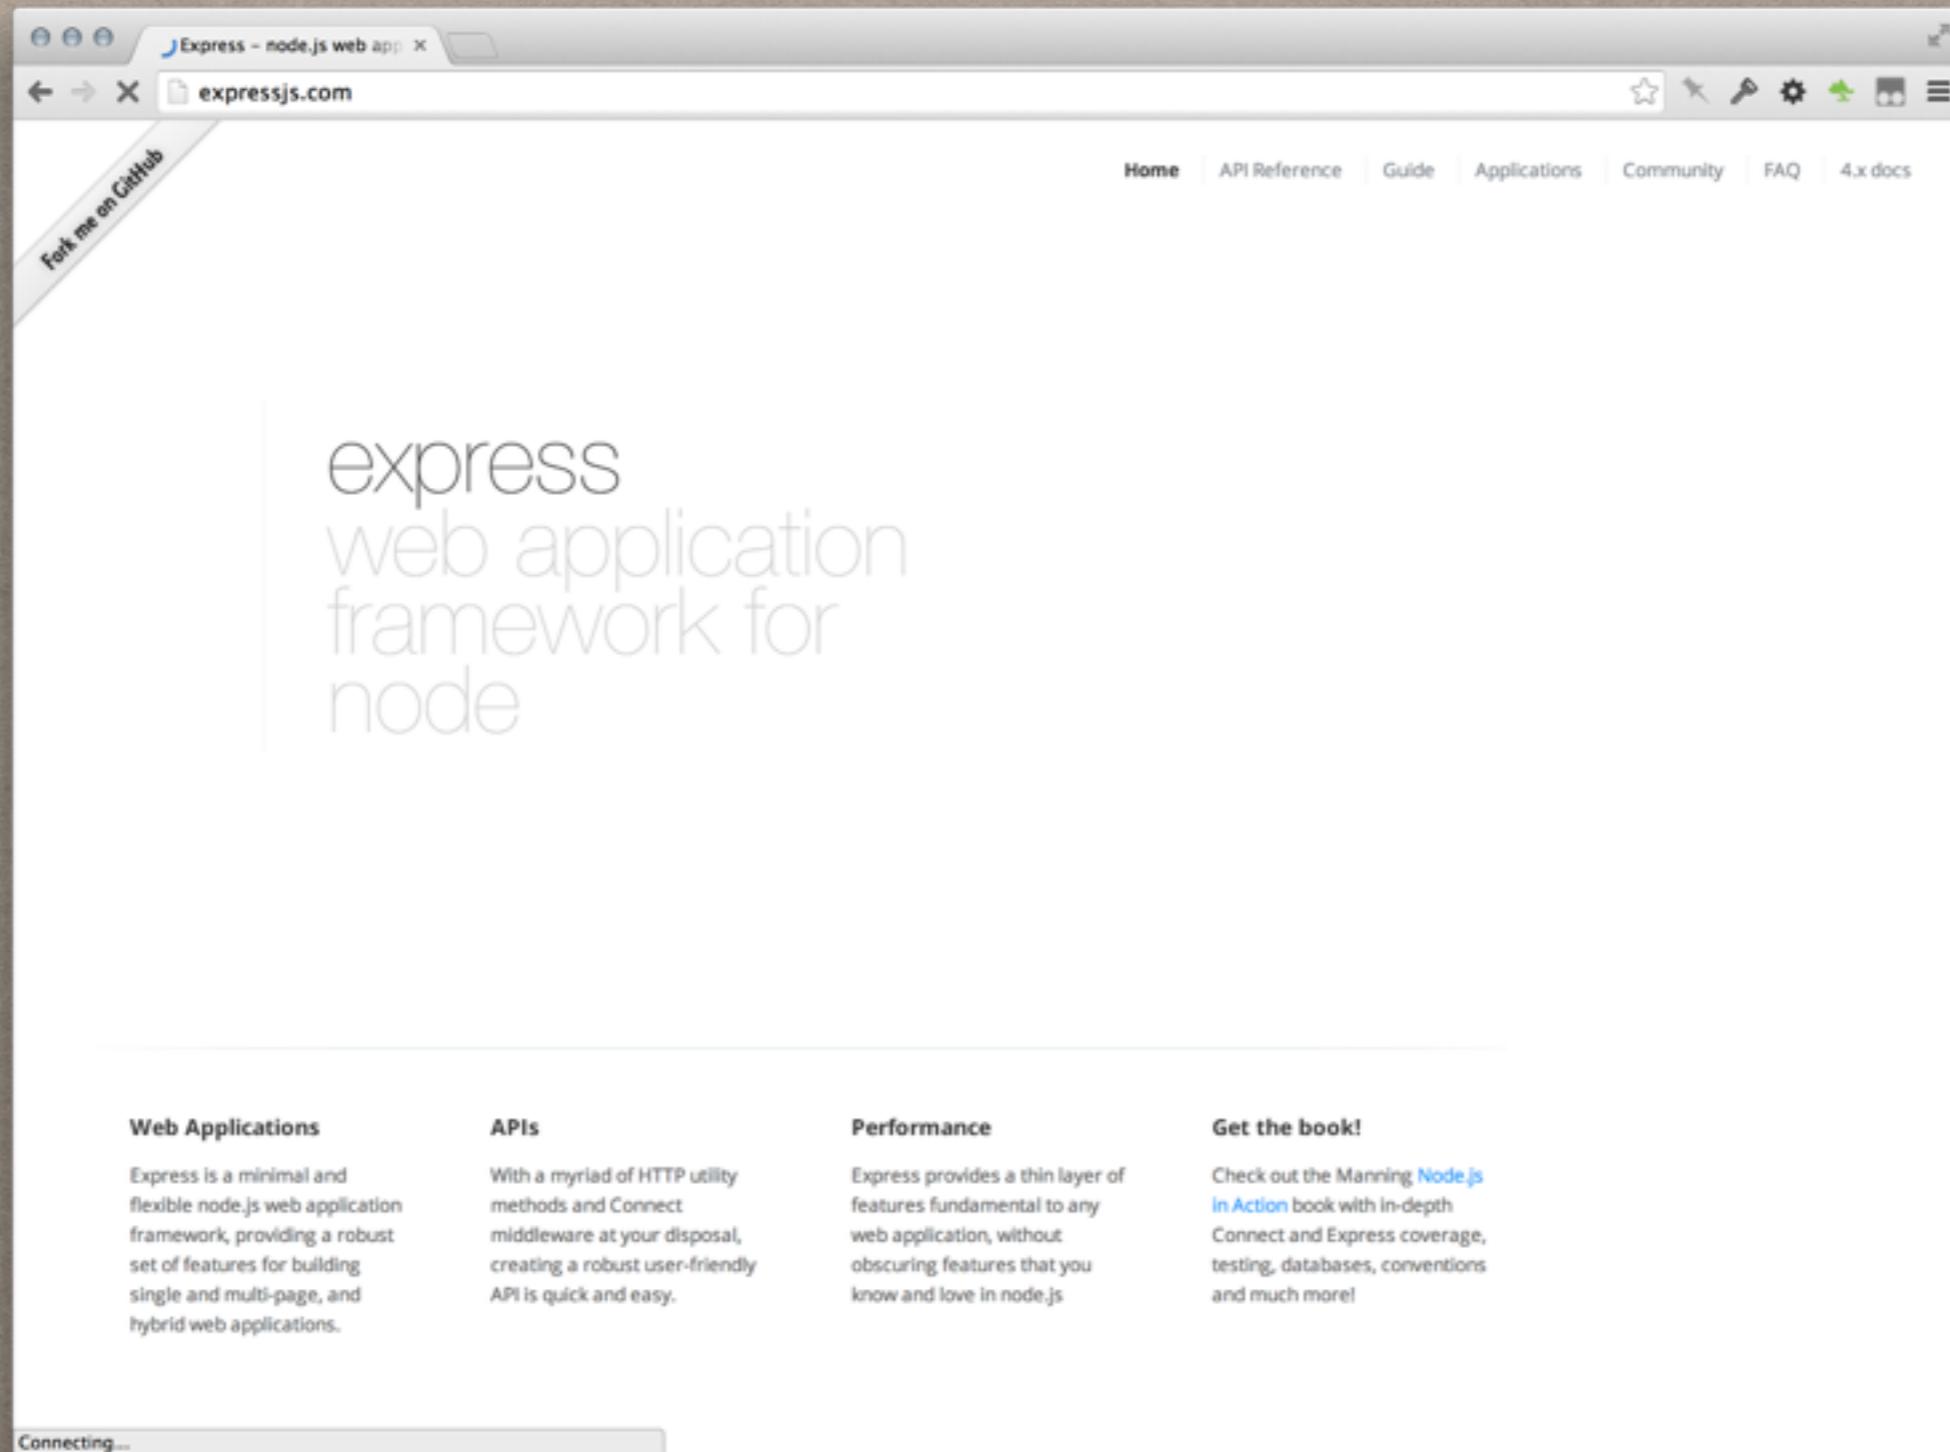
- Logging requests
- Serve “static” files
- Compress responses
- Store posted files
- Parse request parameters and cookies
- Authenticate/
Authorize Requests
- Send non-2xx responses (404, 302, 500, etc.)
- Set cache headers
- Serve dynamic requests

HTTP SERVER FRAMEWORKS

- Connect
- Express
- Restify
- Hapi
- Union

EXPRESS

<http://expressjs.com>

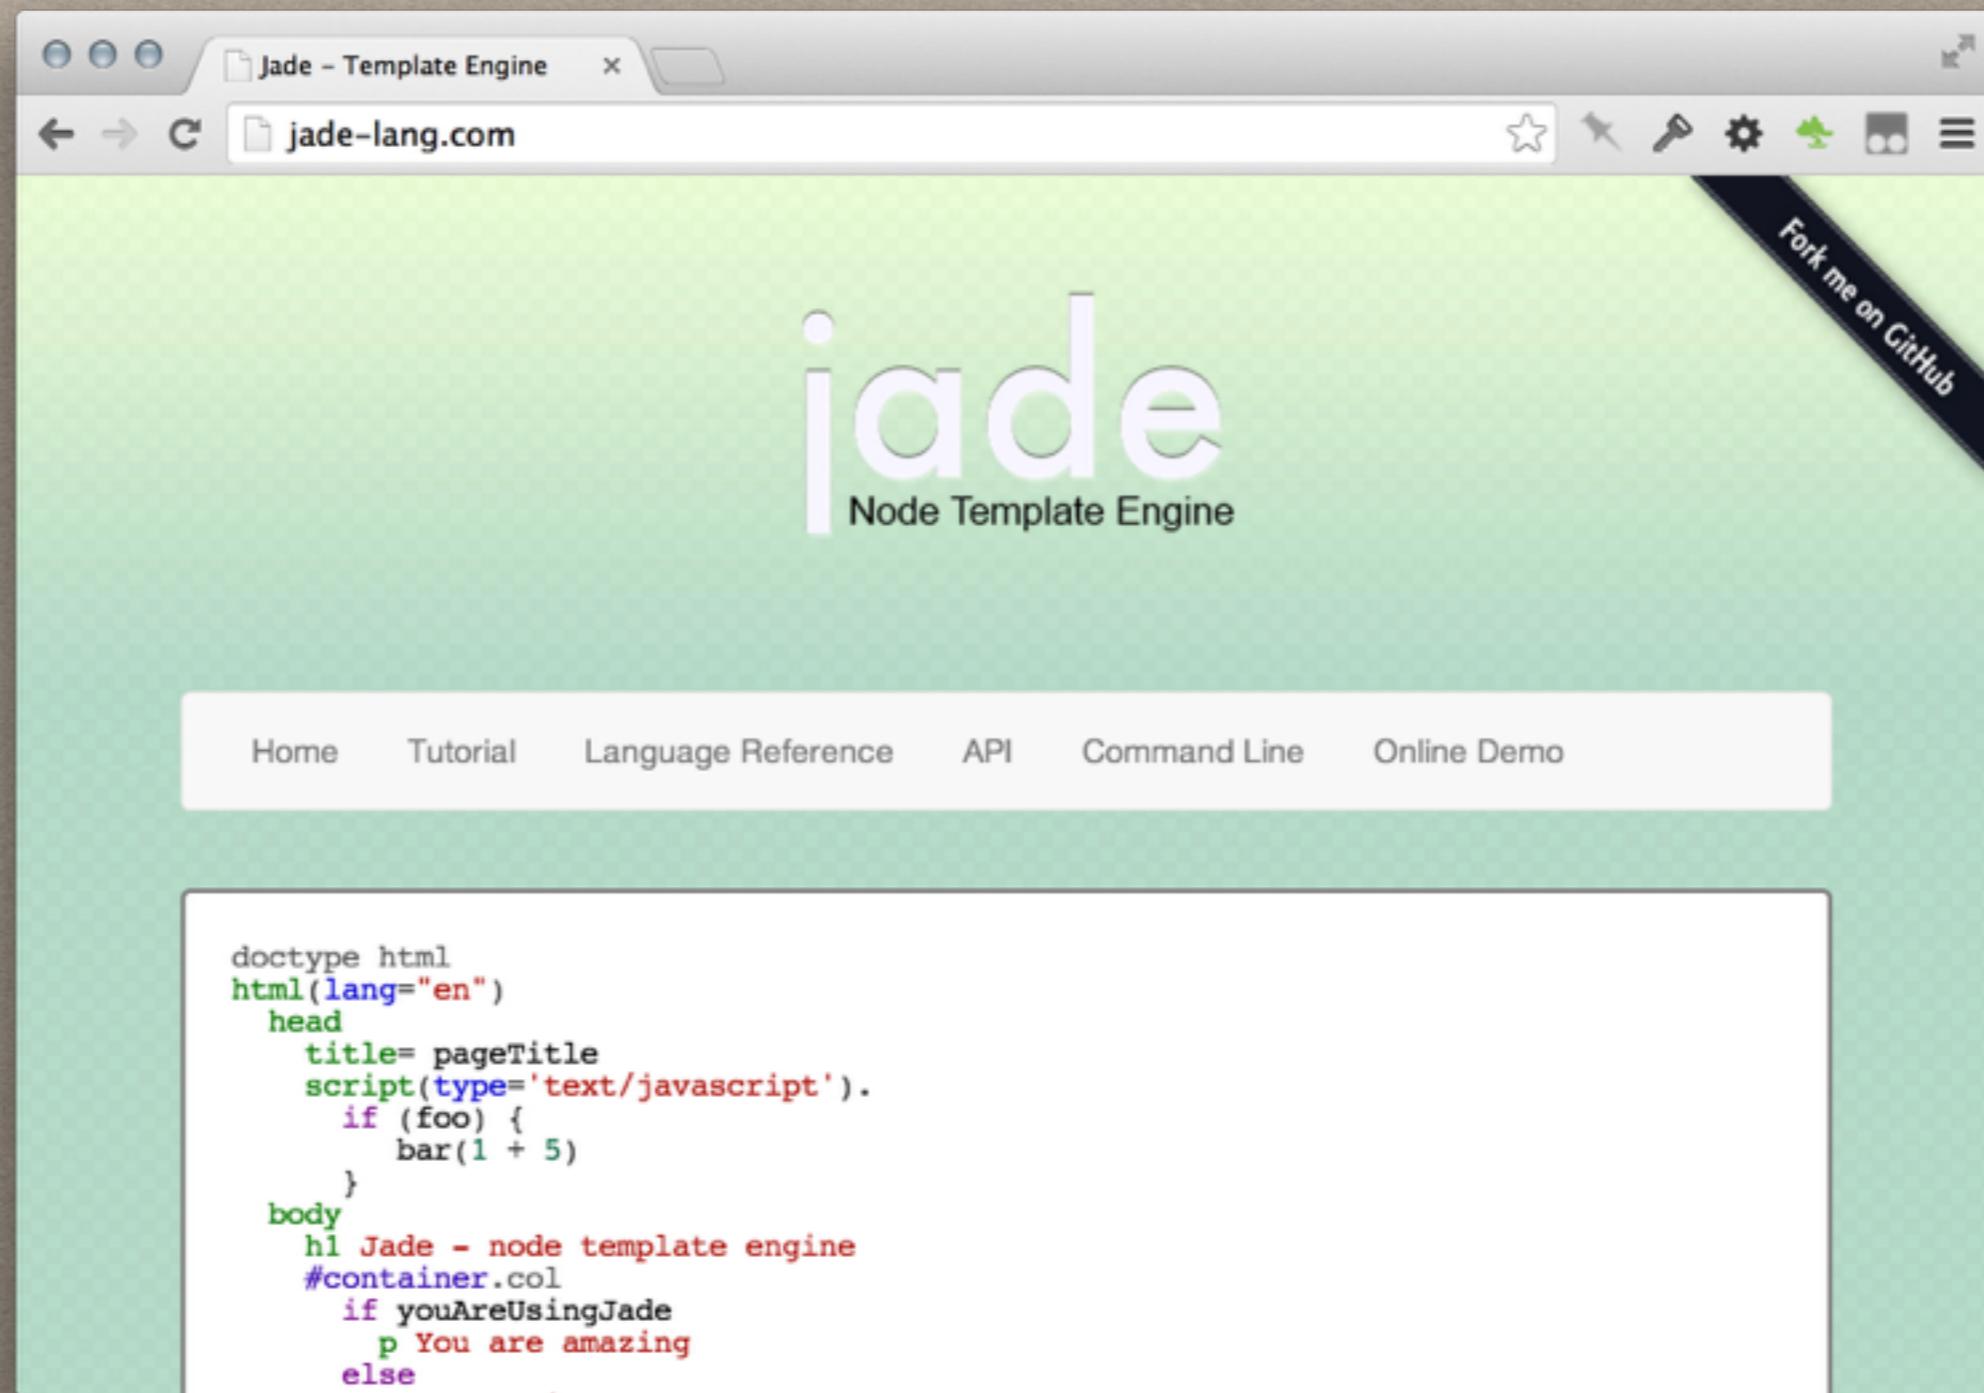


Exercise 5: All Aboard!

- `npm install express`
- Create file “server.js”
- Create express app, create a logging middleware.
- Create a “date time” middleware.
- Leverage express-provided middleware instead of our primitive middleware
- Serve static files from “public” folder.

JADE

<http://jade-lang.com>

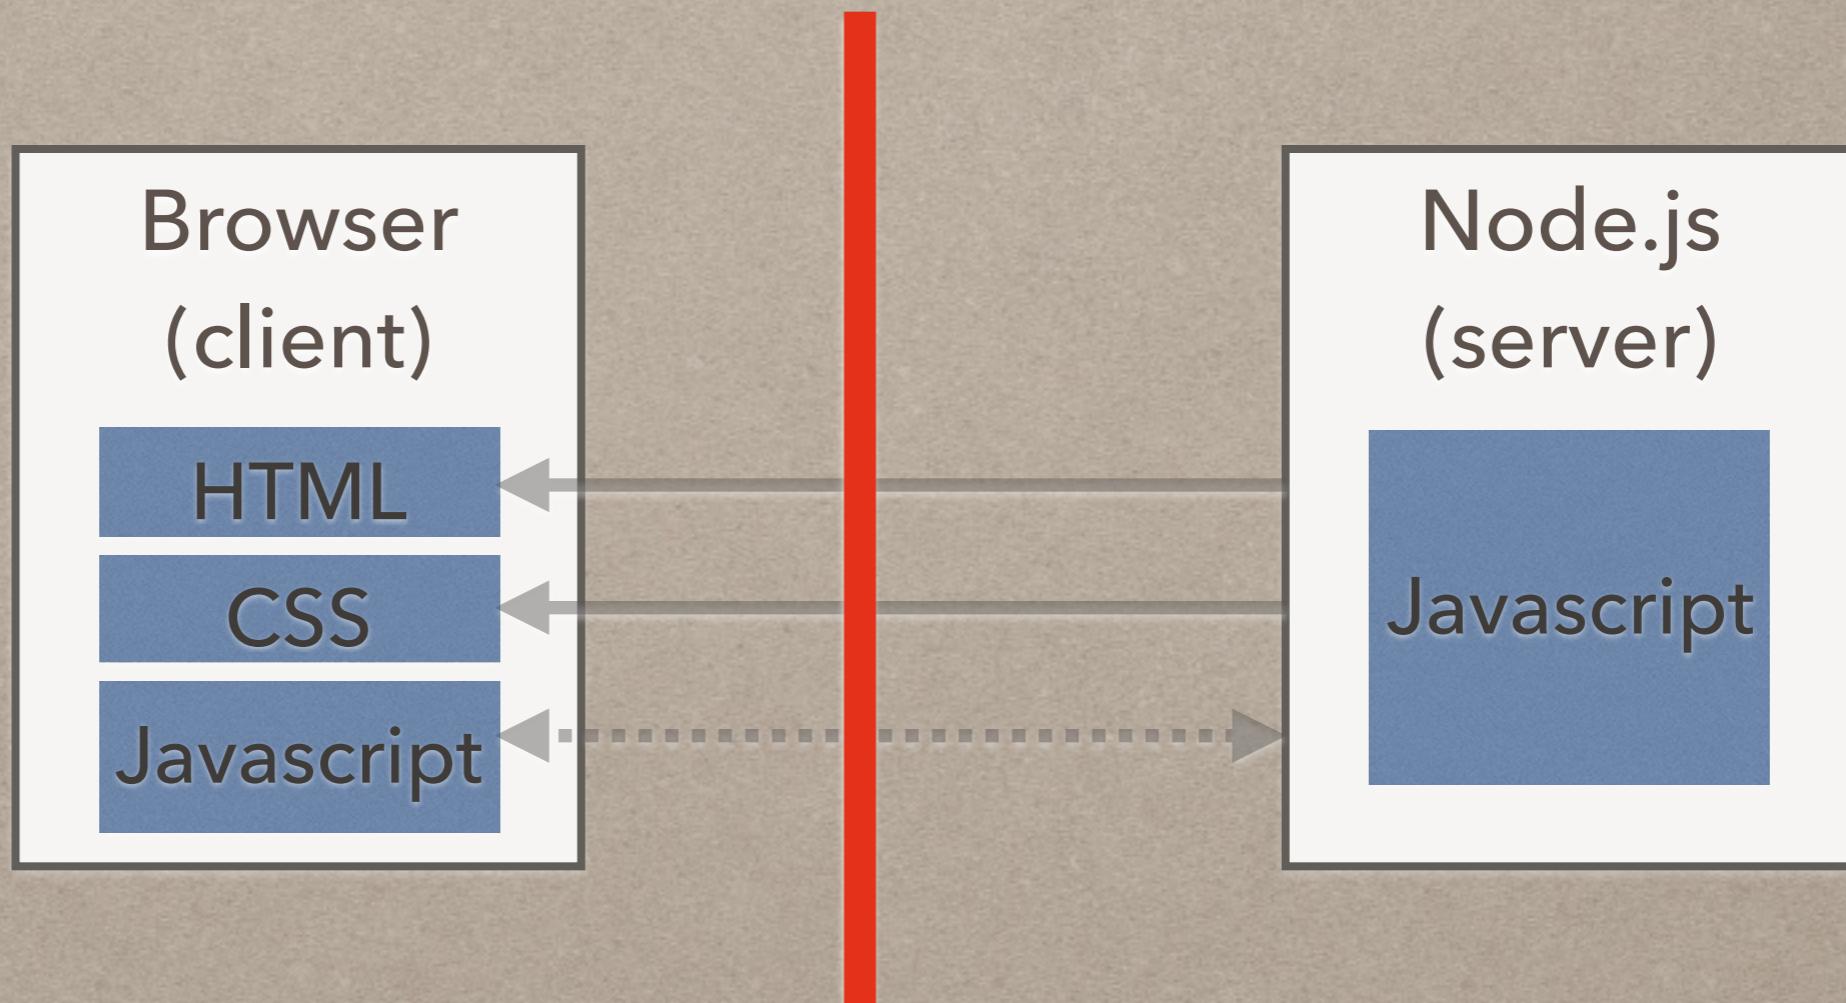


A NOTE ABOUT CLIENT-SERVER COMMUNICATION



← HTTP Request
↔ AJAX Request

A NOTE ABOUT CLIENT-SERVER COMMUNICATION



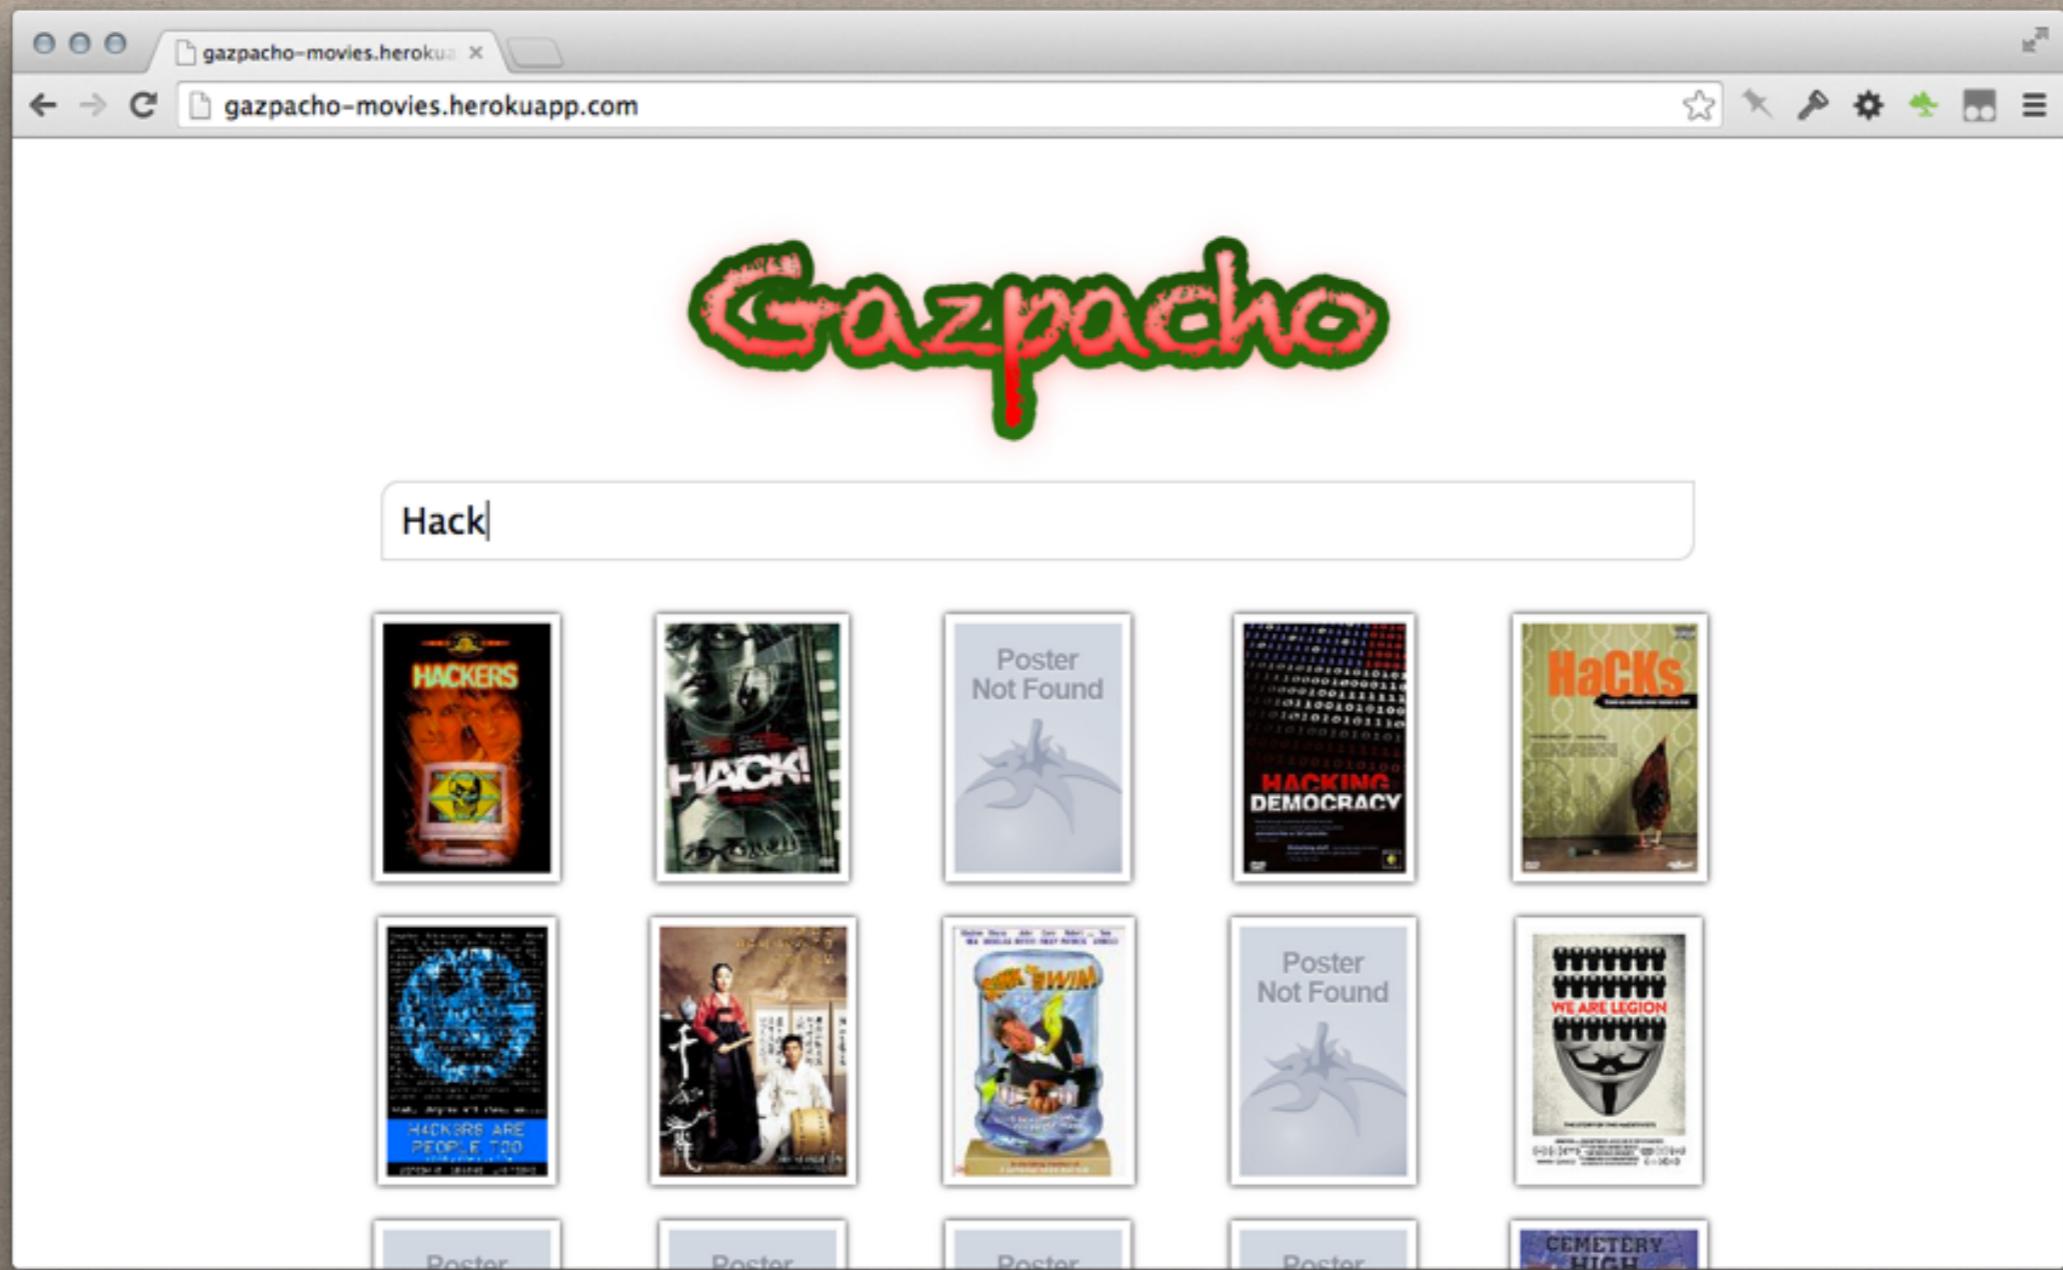
The browser and Node.js
DO NOT share program state.



Exercise 6: I'm Jaded, and you will be, too.

- `npm install jade`
- Configure jade as a rendering engine.
- Create a jade view with a model.

GAZPACHO: A (PRETTY?) BROWSER FOR ROTTEN TOMATOES



Exercise 7: Glue it all together

- `npm install rotten-api`
- Add ‘/listmovies’ API middleware
- Add search box to index.jade
- On keyup, request movies from server.
- Render movie titles as list on client.
- BONUS: Create “favoriting” functionality.

MISCELLANEOUS



PERSISTENCE MODULES

- couchdb
- elasticsearch
- redis
- mongodb
- mysql
- riak
- postgres
- oracle
- voltdb
- aws
- azure
- parse
- csv
- sqlite
- memcached
- cassandra
- sql server

TOOLS

- node-inspector (GUI debugging interface)
- supervisor (Continuously reload node as files are changed)
- “async” and “Seq” (avoid “boomerang” code)
- Grunt (automate repetitive tasks)
- JSHint (Identify syntax errors quickly).
- Bower (easily integrate client-side libraries)
- Yeoman (scaffold new projects)
- Browserify (use some node modules in the browser)
- Jasmine (unit testing)
- git (Source control, very popular in node community)

DEPLOYING

- Heroku (free*)
- Nodjitsu
- Joyent
- Runnable
- c9.io
- Nitrous.io
- Azure
- Web Faction
- Digital Ocean
- Linode
- Rackspace

RESOURCES

- nodeschool.io
- Read [JavaScript: The Good Parts](#)
 - <http://amzn.com/B00260R2ZY>
- Node Philly: <http://node.ph/>
- Node.JS Philly Meetup: <http://www.meetup.com/nodejs-philly/>



Andrew Theken
@atheken
andrewtheke.com



Curtis Herbert